



Department of Computer Science  
UNIVERSITY OF COLORADO **BOULDER**



# Machine Learning: Chenhao Tan

University of Colorado Boulder

LECTURE 20

Slides adapted from Jordan Boyd-Graber, Chris Ketelsen

## Logistics

---

- Homework 4 is due on Sunday!
- Prelim next Friday

## Learning objects

---

- Learn about Adaboost
- Understand the math behind boosting

## Overview

---

Recap of Boosting

Adaboost

The underlying math

## Outline

---

Recap of Boosting

Adaboost

The underlying math

## Boosting intuition

---

Boosting is an ensemble method, but with a different twist.

Idea:

- Build a sequence of dumb models
- Modify training data along the way to focus on difficult to classify examples
- Predict based on weighted majority vote of all the models

Challenges:

- What do we mean by dumb?
- How do we promote difficult examples?
- Which models get more say in vote?

## Boosting intuition

---

What do we mean by dumb?

Each model in our sequence will be a weak learner

$$\text{err} = \frac{1}{m} \sum_{i=1}^m I(y_i \neq h(\mathbf{x}_i)) = \frac{1}{2} - \gamma, \gamma > 0$$

Most common weak learner in Boosting is a decision stump - a decision tree with a single split

## Boosting intuition

---

How do we promote difficult examples?

After each iteration, we'll increase the importance of training examples that we got wrong on the previous iteration and decrease the importance of examples that we got right on the previous iteration

Each example will carry around a weight  $w_i$  that will play into the decision stump and the error estimation

Weights are normalized so they act like a probability distribution

$$\sum_{i=1}^m w_i = 1$$



## Boosting intuition

---

Which models get more say in vote?

The models that performed better on training data get more say in the vote

For our sequence of weak learners:  $h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_K(\mathbf{x})$

Boosted classifier defined by

$$H(\mathbf{x}) = \text{sign} \left[ \sum_{k=1}^K \alpha_k h_k(\mathbf{x}) \right]$$

Weight  $\alpha_k$  is measure of accuracy of  $h_k$  on training data

## Outline

---

Recap of Boosting

Adaboost

The underlying math

## Adaboost

---

- Training set  $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$
- Feature vector  $\mathbf{x} \in \mathbb{R}^d$
- Labels are  $y_i \in \{-1, 1\}$

## Adaboost

---

- 1 . Initialize data weights to  $w_i = \frac{1}{m}, i = 1, \dots, m$
- 2 . For  $k = 1$  to  $K$ :
  - (a) Fit classifier  $h_k(\mathbf{x})$  to training data with weights  $w_i$
  - (b) Compute weighted error  $\text{err}_k = \frac{\sum_{i=1}^m w_i I(y_i \neq h_k(\mathbf{x}_i))}{\sum_{i=1}^m w_i}$
  - (c) Compute  $\alpha_k = \frac{1}{2} \log((1 - \text{err}_k)/\text{err}_k)$
  - (d) Set  $w_i \leftarrow \frac{w_i}{Z_k} \cdot \exp[-\alpha_k y_i h_k(\mathbf{x}_i)], i = 1, \dots, m$
- 3 . Output  $H(\mathbf{x}) = \text{sign} \left[ \sum_{k=1}^K \alpha_k h_k(\mathbf{x}) \right]$

## Adaboost

---

1 . Initialize data weights to  $w_i = \frac{1}{m}, i = 1, \dots, m$

Weights are initialized to uniform distribution. Every training example counts equally on first iteration.

## Adaboost

---

2a. Fit classifier  $h_k(\mathbf{x})$  to training data with weights  $w_i$   
Decide split based on info gain with weighted entropy:

$$H(D) = -p \log_2 p - (1 - p) \log_2 (1 - p)$$

$$I(D, a) = H(D) - \sum_{v \in \text{vals}(a)} \frac{|\{x \in D | x_a = v\}|}{|D|} H(\{x \in D | x_a = v\})$$

$$p = \frac{\sum_{i=1}^m w_i \cdot I(y_i = 1)}{\sum_{i=1}^m w_i \cdot I(y_i = \pm 1)}$$

## Adaboost

---

2b. Compute weighted error

$$\text{err}_k = \frac{\sum_{i=1}^m w_i I(y_i \neq h_k(\mathbf{x}_i))}{\sum_{i=1}^m w_i}$$

## Adaboost

---

2b. Compute weighted error

$$\text{err}_k = \frac{\sum_{i=1}^m w_i I(y_i \neq h_k(\mathbf{x}_i))}{\sum_{i=1}^m w_i}$$

Still gives error rate in  $[0, 1]$




## Adaboost

---

2c. Compute  $\alpha_k = \frac{1}{2} \log((1 - \text{err}_k)/\text{err}_k)$

Models with small err get promoted  
(exponentially)

Models with large err get demoted  
(exponentially)



## Adaboost

---

2d. Set  $w_i \leftarrow \frac{w_i}{Z_k} \cdot \exp[-\alpha_k y_i h_k(\mathbf{x}_i)]$ ,  $i = 1, \dots, m$

If example was misclassified weight goes up

If example was classified correctly weight goes down

How big of a jump depends on accuracy of model

$Z_k$  is just a normalizing constant to ensure that the  $w$ 's are a distribution

## Adaboost

---

3. Output  $H(\mathbf{x}) = \text{sign} \left[ \sum_{k=1}^K \alpha_k h_k(\mathbf{x}) \right]$


Sum up weighted votes from each model

Classify  $y = 1$  if positive and  $y = -1$  if negative

## Adaboost example

---

Suppose we have the following training data





## Adaboost example

---

### First decision stump

---





## Adaboost example

---

First decision stump ,  $err_1 = 0.3, \alpha_1 = 0.42$

---





## Adaboost example

---

### Second decision stump

---





## Adaboost example

---

Second decision stump ,  $err_2 = 0.21$ ,  $\alpha_1 = 0.65$

---







## Adaboost example

---

### Third decision stump

---





## Adaboost example


---

Third decision stump ,  $err_2 = 0.14$ ,  $\alpha_1 = 0.92$

---




## Adaboost example



## Generalization performance

---

Recall the standard experiment of measuring test and training error vs. model complexity




Once overfitting begins, test error goes up

## Generalization performance

---

Boosting has remarkably uncommon effect



Happens much slower with boosting

## Adaboost

---

- 1 . Initialize data weights to  $w_i = \frac{1}{m}, i = 1, \dots, m$
- 2 . For  $k = 1$  to  $K$ :
  - (a) Fit classifier  $h_k(\mathbf{x})$  to training data with weights  $w_i$
  - (b) Compute weighted error  $\text{err}_k = \frac{\sum_{i=1}^m w_i I(y_i \neq h_k(\mathbf{x}_i))}{\sum_{i=1}^m w_i}$
  - (c) Compute  $\alpha_k = \frac{1}{2} \log((1 - \text{err}_k)/\text{err}_k)$
  - (d) Set  $w_i \leftarrow \frac{w_i}{Z_k} \cdot \exp[-\alpha_k y_i h_k(\mathbf{x}_i)], i = 1, \dots, m$
- 3 . Output  $H(\mathbf{x}) = \text{sign} \left[ \sum_{k=1}^K \alpha_k h_k(\mathbf{x}) \right]$

## Quiz

---

Which of the following statement is false?

- A. Adaboost can return the same weak classifier at different  $K$ .
- B. Adaboost can improve the generalization performance with many rounds.
- C. Every misclassified data point has the same multiplicative factor in the update.
- D.  $\alpha_k$  always grows as  $k$  grows.

## Outline

---

Recap of Boosting

Adaboost

The underlying math



## The underlying math

---

So far this looks like a reasonable thing that just worked out  
But is there math behind it?

## The underlying math

---

Yep! It is minimization of a loss function, like always  
Formulate the problem as finding a classifier  $H(\mathbf{x})$  such that

$$H^* = \arg \min_H \sum_{i=1}^m L(y_i, H(\mathbf{x}_i))$$

where here we take the loss function  $L$  to be the following exponential function

$$L = \exp[-y H(\mathbf{x})]$$

Notice if  $y \neq H(\mathbf{x})$  we get a positive exponent and a large loss.

## The underlying math

---

Since we're doing this in an iterative way, we're going to assume a form of  $H(\mathbf{x})$  that is amenable to iterative improvement. Specifically

$$H(\mathbf{x}) = \sum_{k=1}^K \alpha_k \phi_k(\mathbf{x})$$

So the problem becomes to choose the optimal weights,  $\alpha$ , and optimal basis functions  $\phi_k$ .

For everything I will assume that we've already computed a good  $H_{k-1}$  and attempt to build a better  $H_k$ .

## The underlying math

---

At step  $k$  we have

$$L_k = \sum_{i=1}^m \exp [-y_i (H_{k-1}(\mathbf{x}_i) + \alpha \phi(\mathbf{x}_i))]$$

Taking the things we know out of the exponent gives

$$L_k = \sum_{i=1}^m w_{i,k} \exp [-\alpha y_i \phi(\mathbf{x}_i)], \quad w_{i,k} = \exp [-y_i H_{k-1}(\mathbf{x}_i)]$$

Want to choose good  $\alpha$  and  $\phi$  to reduce loss

## The underlying math

---

Can rewrite as

$$L_k = e^{\alpha} \sum_{y_i \neq \phi(\mathbf{x}_i)} w_{i,k} + e^{-\alpha} \sum_{y_i = \phi(\mathbf{x}_i)} w_{i,k}$$

Add zero in a fancy way

$$L_k = (e^{\alpha} - e^{-\alpha}) \sum_{i=1}^m w_{i,k} I(y_i \neq \phi(\mathbf{x}_i)) + e^{-\alpha} \sum_{i=1}^m w_{i,k}$$

## The underlying math

---

Choose  $\phi$  and  $\alpha$  separately.

A good  $\phi$  would be one that minimizes weighted misclassifications

$$h_k = \arg \min_{\phi} w_{i,k} I(y_i \neq \phi(\mathbf{x}_i))$$

That's what our weak learner is for

## The underlying math

---

Plugging that into our Loss function gives

$$L_k = (e^\alpha - e^{-\alpha}) \sum_{i=1}^m w_{i,k} I(y_i \neq h_k(\mathbf{x}_i)) + e^{-\alpha} \sum_{i=1}^m w_{i,k}$$

Now we want to minimize w.r.t.  $\alpha$ . Take derivative, set equal to zero

$$0 = \frac{dL_k}{d\alpha} = (e^\alpha + e^{-\alpha}) \sum_{i=1}^m w_{i,k} I(y_i \neq h_k(\mathbf{x}_i)) - e^{-\alpha} \sum_{i=1}^m w_{i,k}$$

which gives 
$$e^{2\alpha} = \frac{\sum_i w_{i,k}}{\sum_i w_{i,k} I(y_i \neq h_k(\mathbf{x}))} - 1 = \frac{1}{\text{err}_k} - 1 = \frac{1 - \text{err}_k}{\text{err}_k}$$

And finally 
$$\alpha_k = \frac{1}{2} \log \left( \frac{1 - \text{err}_k}{\text{err}_k} \right)$$

## The underlying math

---

What about the weight update?

Remember we got the weights by pulling the already computed function out of  $L$ .  
For the new weights, we have

$$w_{i,k+1} = \exp[-y_i H_k(\mathbf{x}_i)] = \exp[-y_i H_{k-1}(\mathbf{x}_i) - \alpha_k y_i h_k(\mathbf{x}_i)] = w_{i,k} \exp[-\alpha_k y_i h_k(\mathbf{x}_i)]$$

which gives the update

$$w_i \leftarrow w_i \exp[-\alpha_k y_i h_k(\mathbf{x}_i)]$$

And finally,  $H_K(\mathbf{x}) = \text{sign} \left[ \sum_{k=1}^K \alpha_k h_k(\mathbf{x}) \right]$

This is exactly Adaboost